

THE COMMUNICATION PRODUCTIVITY PROBLEM

Your AI copilot that writes like *you*.

Knowledge workers waste 2–3 hours every day crafting professional messages on LinkedIn and Gmail — context-switching, losing contact history, and getting AI replies that sound nothing like them. Comms Assistant is the fix: a browser-native AI that knows your voice, remembers every contact, and drafts replies you'd actually send.

2.8B
TARGET PLATFORM USERS

2–h
DAILY COMMS TIME PER PRO

\$47B
PRODUCTIVITY SOFTWARE TAM

83
TESTS SHIPPED · H1–H5 DONE

Professional communication is *broken*.

For anyone who communicates professionally at scale — recruiters, sales reps, founders, account managers, consultants — the daily reality of LinkedIn and Gmail is an exhausting, context-free grind. Three compounding failures make it so.

28%

of the knowledge worker day is
communication

McKinsey estimates professionals spend 28% of their workday reading and responding to messages — more than any other single activity. Most of that is reactive, context-free, and forgettable.

0

existing AI tools remember who
you're talking to

ChatGPT, Copilot, Gemini — they all generate from scratch. No memory of your last conversation with this person. No record of what you promised. No sense of relationship history. Every message starts cold.

94%

of users say AI-drafted messages
"don't sound like me"

Generic AI outputs destroy the authenticity that makes professional relationships work. People can tell when a message wasn't written by the sender — and it erodes trust at exactly the moment you're trying to build it.

“

*"The problem isn't that AI can't write. The problem is that AI doesn't know **who you are** — your voice, your relationships, your history with every person you've ever messaged."*

— CORE INSIGHT BEHIND COMMS ASSISTANT

A communication memory layer for the *professional web*.

Comms Assistant is a Chrome extension that lives inside LinkedIn and Gmail — where the work already happens. It does three things no existing tool does together: it learns your writing voice, it remembers every professional relationship, and it drafts context-aware replies that sound exactly like you.

1

It reads the conversation

The extension's content script extracts the live thread from LinkedIn or Gmail — participant names, message history, tone, intent — without the user copying or pasting anything.

2

It knows who you're talking to

A persistent SQLite memory layer stores everything known about this contact: past interactions, confirmed notes, relationship strategies, follow-up commitments. Context that would otherwise live in your head — or be forgotten.

3

It drafts in your voice

A one-time voice profile distillation — trained on your own past messages — gives the LLM a precise style guide. Not a generic tone selector. Your actual patterns, vocabulary, and rhythm.

4

It learns over time

Every confirmed note, edited draft, and piece of feedback refines the system. The assistant gets sharper with each interaction — not static, always improving.



Why this is *hard to copy*.

The surface product is a Chrome extension. The real asset is three compounding data advantages that make the product better the longer you use it — and significantly harder to replicate the longer we have the market.



Voice Profile Flywheel

The voice profile isn't a tone selector — it's a statistical model of your writing distilled from your actual corpus. Every confirmation, edit, and rejection makes it more accurate. Competitors starting from scratch with a new user face months of flywheel disadvantage.

DATA COMPOUNDING



Relationship Memory Graph

Contact notes, strategies, and follow-up context accumulate into a relationship intelligence layer. Switching costs are inherently high — your contact memory doesn't export to ChatGPT. The longer you use it, the more you'd lose by leaving.

SWITCHING COST



Privacy-First Architecture

Designed for zero-trust from day one: all data is tenant-scoped, API keys are AES-256-GCM encrypted at rest, and the extension can run fully local. Enterprise buyers who've been burned by SaaS data leaks will pay a premium for this.

ENTERPRISE READY



Platform-Agnostic Design

LinkedIn and Gmail are live today via a pluggable PlatformExtractor interface. WhatsApp Web, X DMs, Outlook — any platform that runs in a browser is an addable platform. Distribution expands without a rewrite.

EXTENSIBLE



Context-Injected at the Prompt Level

Prompt injection attacks are a real risk for any AI that reads user communication. We designed the trust boundary in from the start: conversation content is explicitly tagged as untrusted data, separate from the user's own voice and confirmed memory.

SECURITY MOAT



Multi-Tenant Infrastructure

Full multi-tenancy is live — bearer-token auth, per-tenant isolation, BYO LLM keys, rate limiting, and data export/purge. Going from a single-user local tool to a hosted SaaS requires a full rewrite for most competitors. We already did it.

INFRA MOAT

A massive market hiding in plain *sight*.

The professional communication software market is large, growing, and underserved by AI that actually knows you.

\$47B	Total Addressable Market Global productivity software, 2026
1B+	LinkedIn Users Professional network, primary platform
1.8B+	Gmail Active Users Second platform, live and expanding
\$8B	Serviceable Obtainable Market High-volume professional communicators

The Ideal Customer is Already in Pain

Recruiters (300M on LinkedIn), sales professionals, founders doing outbound, investor relations teams, enterprise account managers — these users send 50–200 professional messages per day. Even a 30-minute daily time saving at \$75/hr is worth \$150/month in productivity. The value proposition is obvious.

Browser Extensions Are Underrated Distribution

The Chrome Web Store is a discovery engine with 3B+ Chrome users. A well-positioned extension gets organic installs from people searching for exactly this solution. Grammarly grew to 30M daily users almost entirely through this channel. We're building on the same rails.

Enterprise Motion Is Already Designed In

The multi-tenant architecture, Bearer-token auth, BYO LLM keys, and data portability (GDPR-compliant export/purge) aren't features we'll add later — they're live and tested. The technical proof that this scales to enterprise accounts exists today.

Three tiers. Clear *upgrade logic*.

Freemium drives top-of-funnel at zero CAC. Per-seat captures teams. Enterprise BYO-key model solves the price objection for large deployments.

FREE \$0 forever · local only Solo Professional - Full extension on LinkedIn + Gmail - Local SQLite memory — fully private - Voice profile from your own corpus - Gemini CLI provider (free tier) - No backend required	RECOMMENDED \$19 per seat / month · hosted Team - Hosted backend on managed infra - Per-member tenant isolation - Any OpenAI-compatible LLM - Shared admin + provisioning CLI - Rate-limited, always-on backend - GDPR export + purge per user	ENTERPRISE Custom contract · on-prem or hosted Enterprise - BYO LLM API key (AES-256-GCM encrypted) - On-prem Docker or dedicated hosted - SSO + enterprise auth integration - Volume rate limits + SLA - Data residency guarantees - Custom platform extractors
--	---	---

Not a deck. *A working product.*

Built, tested, and deployed. The full product — from Chrome extension to hosted multi-tenant backend — exists today and is production-ready.



Chrome Extension — Live

MV3 extension with Shadow DOM overlay on LinkedIn and Gmail. Platform-agnostic extractor interface. Popup for backend config. Packaged and ready for Chrome Web Store submission.

MV3 · TypeScript · React · Vite



Production Backend — Deployed

Node/Express backend with full multi-tenant support. SQLite WAL with composite PK/FK schema. Bearer-token auth, rate limiting, AES-256-GCM secrets. Dockerfile ready for Fly.io or Railway.

Docker · SQLite WAL · Express ESM



LLM Layer — Pluggable

Two live providers: Gemini CLI for local use and an OpenAI-compatible HTTP adapter for hosted. Per-tenant provider routing. Voice quality eval harness with retry logic. BYO key support.

Gemini · OpenAI-compatible · Per-tenant



Security — Enterprise Grade

SHA-256 bearer token storage (never plaintext), AES-256-GCM secrets at rest with scrypt key derivation, per-tenant filesystem isolation with path traversal guards, explicit prompt injection trust boundary.

AES-256-GCM · SHA-256 · scrypt



Test Suite — 83 Tests

Vitest test suite covering DB migrations, auth, encryption, rate limiting, tenant isolation, and LLM prompt generation. GitHub Actions CI runs on every push. 14 test files, all passing.

83 tests · 14 files · CI green



Dashboard — React Console

Full contact management dashboard with follow-up calendar (.ics export), voice analysis view, activity feed, live LLM settings, hash routing, and deep-link contact views.

React · dark/light theme · hash routing

Built to scale.

Engineered correctly.

The technical decisions weren't made for a demo. They were made for a product that will serve millions of users on shared infrastructure with enterprise compliance requirements.

// MULTI-TENANCY		// SECURITY	
DB isolation	tenant_id on all tables	Token format	cmst_<32B base64url>
Contacts PK	(tenant_id, name) composite	Token storage	SHA-256 hash only
Note FK	REFERENCES contacts ON CASCADE	API key enc	AES-256-GCM + scrypt
FS isolation	data/tenants/<id>/ per tenant	Prompt boundary	<UNTRUSTED_CONVERSATION> fence
Path guard	safeTenant() sanitizer	CORS	Configurable allowlist
Migration guard	PRAGMA user_version	Rate limit	Sliding-window, Redis-swappable
// EXTENSION		// DEPLOYMENT	
Manifest	MV3 (future-proof)	Container	node:20-slim (glibc)
UI injection	Shadow DOM (no style bleed)	Persistence	/app/data volume
Backend URL	chrome.storage.sync	DB mode	WAL (write-ahead log)
Host perms	optional (user-granted)	Platforms	Fly.io · Railway · on-prem
Platform iface	PlatformExtractor (pluggable)	Provisioning	tenant CLI via docker exec
Build	Vite + @crxjs	CI	GitHub Actions

Builders who shipped, not *presenters who promised.*

Two founders. The full product — extension, backend, multi-tenancy, security, CI — was built and shipped by this team.
No outsourcing. No technical debt shortcuts.



Vaidurya Shah
FOUNDER · FULL-STACK

Designed and built the core product architecture — Chrome extension (MV3), the Node/Express backend, multi-tenant data model, bearer-token auth, and the LLM routing layer. Led the H1–H5 hosting phases from concept to production deployment.

 `Vaidurya-s`



Saral Gupta
CO-FOUNDER · ENGINEERING

Owns the security architecture, encryption layer (AES-256-GCM, scrypt), test suite (83 tests, GitHub Actions CI), and platform abstraction system. Implemented the Gmail extractor, snapshot workflow, and tenant data lifecycle (export/purge).

 `saralgupta09`

WHERE WE'RE GOING

*"Every professional relationship you'll ever have starts with a message. We're building the infrastructure that makes **every one of those messages** your best."*

LinkedIn and Gmail are phase one. The long-term vision is a universal communication memory layer — a personal relationship OS that travels with you across every professional platform, knows who you are, and helps you show up at your best in every interaction. Not just AI. Your AI.